

Bloomberg

November 18, 2024

1 Text Analysis of Bloomberg Articles

1.1 Analyzing company sentiment using NLP models, Data Cleaning and EDA

1.2 Regex, Pandas, Seaborn, Neural Network Models, Data Pipelines

Working with Bloomberg news articles on Microsoft and Microsoft stock data (MSFT).

Summary of methods:

- Conducting data cleaning and EDA on a text-based dataset,
- Manipulating data in `pandas` with the `datetime` and `string` accessors,
- Writing regular expressions and using `pandas` RegEx methods, and
- Performing sentiment analysis on text using DistilBERT.

```
[37]: import warnings
warnings.simplefilter(action="ignore")

import re
import itertools
import numpy as np
import pandas as pd
import seaborn as sns
import matplotlib.pyplot as plt

from utils import *

# Ensure that pandas shows at least 280 characters in columns, so we can see
# ↪ full articles.
pd.set_option("max_colwidth", 280)
plt.style.use("fivethirtyeight")
sns.set()
sns.set_context("talk")
```

In this assignment, we will use the DistilBERT model, which is a Natural Language Processing (NLP) model designed to understand human language by processing text to capture the context and meaning of words within sentences.

```
[9]: from transformers import pipeline
model_checkpoint = "distilbert/distilbert-base-uncased-finetuned-sst-2-english"
```

1.3 Importing the Data

The data for this assignment is a subset of the financial news dataset from [this github repo](#). The original datasets are no longer available online due to copyright issues, but I was allowed access for educational purposes. The data in the file `data/msft_bloomberg_news.txt` has been filtered to just Bloomberg articles published between 2010 to 2013 (inclusive) with text that contains “Microsoft” or “MSFT” (Microsoft’s stock name).

1.3.1 EDA

```
[10]: file = 'data/msft_bloomberg_news.txt'
      with open(file) as f:
          q1a = f.read(1000)
      print(q1a)
```

```
[{"id":46243185,"title":"Opera Jumps Most Ever After Report Facebook May Bid:
Oslo Mover","released_at":"<date>May 29 2012</date>
<time>09:40:58</time>","content":"Opera Software ASA (OPERA) , the
Norwegian\marker of Internet browsers, surged the most on record in Oslo\after
technology website Pocket-Lint reported that Facebook Inc. (FB) \nmay try to
acquire the company. Opera gained as much as 26 percent, the biggest jump
since\nit first sold shares in 2004. The Oslo-based company rose 18\npercent to
40.5 kroner at 11:37 a.m., giving it a market value\nof 4.85 billion kroner
($807 million). Opera is the last major independent browser left, with
the\nothers owned by companies such as Microsoft Corp. (MSFT) , Google Inc.
(GOOG) \nand Apple Inc. (AAPL) , said Aleksander Nilsen, an analyst at Abg
Sundal\nCollier in Oslo. The company has a strong balance sheet, and\ncould be
an attractive target for other companies, such as\nMountain View , California-
based Google, he said.
```

```
[11]: q1b = 'C'
```

```
[12]: grader.check("q1b")
```

```
-----
NameError                                Traceback (most recent call last)
Cell In[12], line 1
----> 1 grader.check("q1b")

NameError: name 'grader' is not defined
```

```
[13]: msft_news_df = pd.read_json(r'data/msft_bloomberg_news.txt').set_index('id')
      msft_news_df.head(1)
```

```
[13]:                                     title \
id
```

```

46243185  Opera Jumps Most Ever After Report Facebook May Bid: Oslo Mover

                                released_at  \
id
46243185  <date>May 29 2012</date> <time>09:40:58</time>

                                content  \
id
46243185  Opera Software ASA (OPERA) , the Norwegian\marker of Internet
browsers, surged the most on record in Oslo\after technology website Pocket-
Lint reported that Facebook Inc. (FB) \may try to acquire the company. Opera
gained as much as 26 percent, the biggest jump since\n...

                                path
id
46243185  ./2008_2012_msft_bloomberg_news/opera-jumps-most-on-record-after-
report-of-facebook-s-interes.txt

```

1.3.2 Analysis

Specating if we are interested in using the news to predict future stock values. What additional data would we need to predict stock prices

To assist in predicting future stock values it would be helpful to merge columns that join the corresponding company's stock ticker based on the time under 'released_at'. Additionally I would parse through the title and content to identify key words that may allude to an increase in stock valuation of a company such as "jumps" or "surges" (wording used in first entry columns).

1.4 Time Analysis

After loading in the data, we can start exploring news articles by analyzing the relationships between the release dates (date of publication) and different topics and companies.

EXAMPLES: of time

```

<date>May 29 2012</date> <time>09:40:58</time>
<date>May 18 2011</date> <time>22:42:40</time>
<date>August 15 2012</date> <time>00:09:02</time>
<date>July 1 2011</date> <time>22:12:37</time>
...

```

There are several ways to convert this to a `Timestamp` object that we can use more easily. Use regex such that: 1. The same index as `msft_news_df` (id) and 2. Column labels: "Month", "Day", "Year", "Hour", "Minute", "Second".

convert all numerical values ("Year", "Day", "Hour", "Minute", "Second") to type `int`.

```

[14]: time_names = ["Month", "Day", "Year", "Hour", "Minute", "Second"]
      data = pd.DataFrame(columns = time_names)

```

```
[15]: pattern = '>{1}([A-Za-z+)]\s+(\d{1,2})\s(\d{4})\s</\s+[date]+\s+\s<[time>]+(\d+):
      ↪(\d+):(\d+)<'
      dates = msft_news_df["released_at"].str.extract(pattern)
      dates
      for i in np.arange(1,6):
          dates[i] = dates[i].astype(int)
      dates = dates.rename({0:"Month", 1:"Day", 2:"Year", 3:"Hour", 4:"Minute", 5:
      ↪"Second"}, axis = 1)
      dates
```

```
[15]:
```

	Month	Day	Year	Hour	Minute	Second
id						
46243185	May	29	2012	9	40	58
73522879	May	18	2011	22	42	40
29296500	August	15	2012	0	9	2
49799724	July	1	2011	22	12	37
20739032	January	18	2012	1	20	28
...
75325873	June	27	2012	0	35	58
49071474	September	24	2013	13	38	57
12417018	September	14	2011	4	1	0
25935811	June	28	2010	1	0	0
21143940	September	8	2011	1	11	1

[4635 rows x 6 columns]

```
[16]: temp = pd.merge(msft_news_df, dates, left_index = True, right_index =
      ↪True)[["title", "released_at", "content", "path", "Month", "Day", "Year"]]
      msft_news_2010 = temp.loc[temp['Year'] ==2010]
      msft_news_2010.head(1)
```

```
[16]:
```

	title \
id	
95357231	Netflix Profit Jumps 44% on New Users

	released_at \
id	
95357231	<date>April 21 2010</date> <time>23:52:36</time>

	content \
id	
95357231	Netflix Inc. said first-quarter\nprofit rose 44 percent as the movie subscription service signed\nup new customers and increased online offerings. \n Net income advanced to \$32.3 million, or 59 cents a share,\nfrom \$22.4 million, or 37 cents, a year earlier, the Los Gatos,\nn...

	path \
--	--------

```
id
95357231  ./2008_2012_msft_bloomberg_news/netflix-quarterly-profit-
increases-44-as-movie-rental-service-adds-users.txt
```

```
Month Day Year
id
95357231 April 21 2010
```

adding a boolean column to the `msft_news_df` DataFrame indicating whether the corresponding company is mentioned in the text of the article. Also six new columns containing True/False values to the DataFrame: "amazon", "nintendo", "apple", "sony", "facebook", "netflix".

```
[17]: companies = ["amazon", "nintendo", "apple", "sony", "facebook", "netflix"]

#for loop over each row in msft df and run another for loop iterating through
↳ each string in companies
#secondly assign each new company column a boolean whether the respective
↳ company is contained in the content
#str.lower() the content column
msft_news_df['content'] = msft_news_df['content'].str.lower()
for company in companies:
    msft_news_df[company] = msft_news_df['content'].str.contains(company, case
↳ = True)
msft_news_df.head(1)
```

```
[17]: title \

id
46243185 Opera Jumps Most Ever After Report Facebook May Bid: Oslo Mover

released_at \

id
46243185 <date>May 29 2012</date> <time>09:40:58</time>

content \

id
46243185 opera software asa (opera) , the norwegian\marker of internet
browsers, surged the most on record in oslo\after technology website pocket-
lint reported that facebook inc. (fb) \may try to acquire the company. opera
gained as much as 26 percent, the biggest jump since\n...

path \

id
46243185  ./2008_2012_msft_bloomberg_news/opera-jumps-most-on-record-after-
report-of-facebook-s-interes.txt

amazon nintendo apple sony facebook netflix
id
```

46243185 False False True False True False

```
[18]: companies + ['Year']
```

```
[18]: ['amazon', 'nintendo', 'apple', 'sony', 'facebook', 'netflix', 'Year']
```

```
[19]: merged = pd.merge(msft_news_df, dates, left_index = True, right_index = True).
      ↪loc[:, companies + ['Year']]
      year_news = merged.loc[merged['Year'] >= 2010].groupby('Year').sum()
      year_news
```

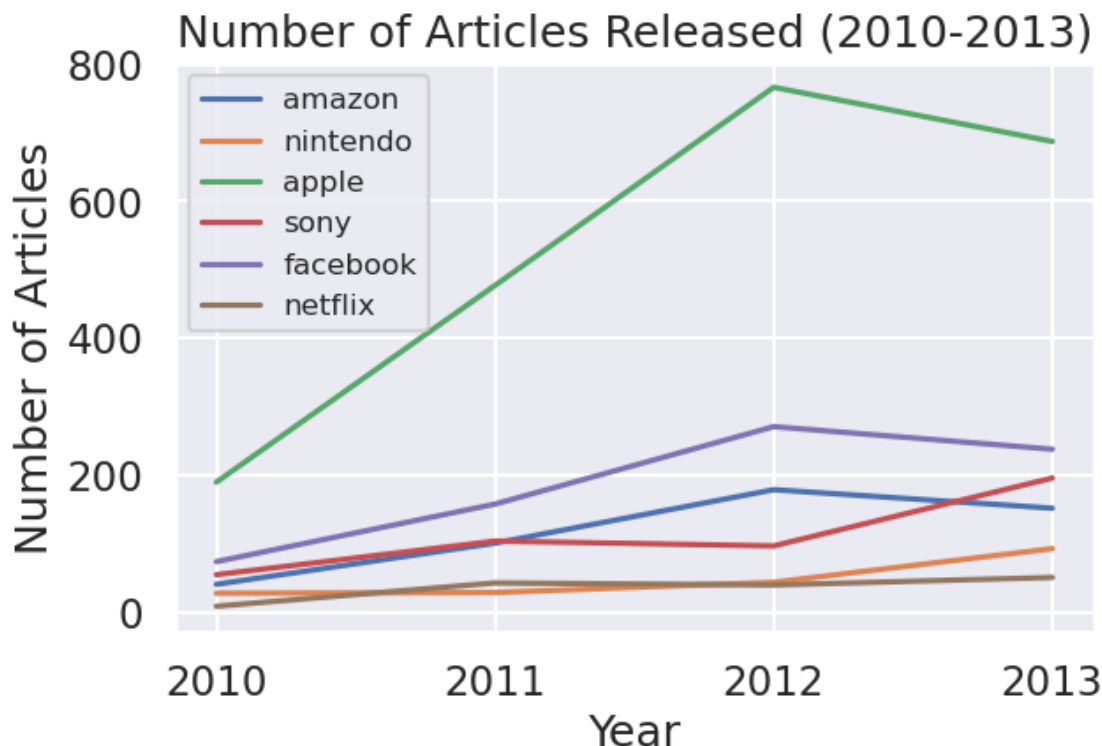
```
[19]:
```

	amazon	nintendo	apple	sony	facebook	netflix
Year						
2010	41	28	190	55	74	9
2011	101	29	477	104	158	43
2012	179	44	766	97	271	40
2013	152	93	687	196	238	51

Plot # of articles

```
[20]: plt.figure(figsize=(6, 4))

      for company in companies:
          sns.lineplot(data=year_news.reset_index(),
                        x="Year",
                        y=company,
                        label=company)
      plt.legend(fontsize="12")
      plt.xticks(np.arange(2010, 2014), np.arange(2010, 2014))
      plt.ylabel("Number of Articles")
      plt.xlabel("Year")
      plt.title("Number of Articles Released (2010-2013)");
```



Reflection:

According to the plot, Apple received the most mentions within bloomberg articles from 2010 - 2013 relative to the other major companies by a significant margin. Pointing to the peak in 2012, there can be several factors and events in Apple's timeline that may have influenced this and caused media coverage. Such events includes the loss in life of Steve Jobs in 2011, launch of the iPhone 5, trial case versus the rival Samsung and morality issues surrounding the outsourcing to China. As with any tragic loss of a public figure representing a major company or a pivotal court case, you can expect headlines surrounding these topics from major journals.

Meanwhile, other companies tend to experience steady and consistent mentions throughout the plot. The data does not capture a company's relative performance as these mentions in the articles can have a binary positive or negative story. Thus, the data without further manipulation cannot entail the performance of different companies as it is restricted to capturing the bloomberg coverage which can be biased based on the business model of Bloomberg and who they choose to report on.

1.5 Sentiment Analysis

We will use a fine-tuned version of the **DistilBERT** model ([github](#), [original paper](#)) to analyze the sentiment of Bloomberg news articles. DistilBERT is a neural network-based language model (a close relative to ChatGPT);

We can use the [HuggingFace library](#) to build the sentiment analysis pipeline and load the model.

```
[21]: # Load the model
sentiment_analysis = pipeline("sentiment-analysis", model=model_checkpoint)

# Get the sentiment of a given string
sentiment_1 = sentiment_analysis("I have two dogs.")
print("Example 1: " + str(sentiment_1))

sentiment_2 = sentiment_analysis("I do not have dogs.")
print("Example 2: " + str(sentiment_2))

sentiment_3 = sentiment_analysis("Fortunately, I do not have dogs to worry_
↳about.")
print("Example 3: " + str(sentiment_3))
```

```
Example 1: [{'label': 'POSITIVE', 'score': 0.9955033659934998}]
Example 2: [{'label': 'NEGATIVE', 'score': 0.9987561702728271}]
Example 3: [{'label': 'POSITIVE', 'score': 0.9975079298019409}]
```

As you can see, the model can determine the sentiment of phrases/sentences (not just words). The model measures the phrase's **polarity**, indicating how strongly negative or positive it is on a scale of 0 to 1.

```
[22]: sentiments = sentiment_analysis(["I have two dogs.", "I do not have dogs."])
print(sentiments[1])
```

```
{'label': 'NEGATIVE', 'score': 0.9987561702728271}
```

1.5.1 To note

Due to model input size constraints, a maximum of 512 words (tokens), and limited computational resources on Datahub, we cannot load the full articles into the pipeline. Instead, we can look at the first sentence that mentions Microsoft in each article.

```
[23]: microsoft_re = '([^.!]*(?:microsoft|msft)+[^.!]*[.?!]+)'
msft_news_2010['content'] = msft_news_2010['content'].str.lower()
msft_news_2010['first_sentence'] = msft_news_2010['content'].str.
↳findall(microsoft_re).str[0]
msft_news_2010.head(5)
```

```
[23]:
```

id	title \
95357231	Netflix Profit Jumps 44% on New Users
75227517	Republican Win May Be Tax Boon for Companies, High Incomes
57850804	Alibaba Says It Now Offers Sohu's Search Engine
75532360	Slim Solution for Trade Imbalances Is More Buying by China
10176588	S&P 500 to Defy 'New Normal' and Rally 17%, Cambiar's Barish Says

	released_at \
id	
95357231	<date>April 21 2010</date> <time>23:52:36</time>
75227517	<date>November 3 2010</date> <time>16:46:00</time>
57850804	<date>October 29 2010</date> <time>12:23:43</time>
75532360	<date>October 31 2010</date> <time>16:05:40</time>
10176588	<date>December 1 2010</date> <time>20:38:58</time>

	content \
id	
95357231	netflix inc. said first-quarter\nprofit rose 44 percent as the movie subscription service signed\nup new customers and increased online offerings. \n net income advanced to \$32.3 million, or 59 cents a share,\nfrom \$22.4 million, or 37 cents, a year earlier, the los gatos,\nn...
75227517	americans with the highest incomes\nand u.s. corporations, especially those with international\noperations, stand to be big winners as newly elected\ncongressional republicans signal they will extend existing tax\nbenefits and push for new ones. republicans will use their ne...
57850804	alibaba group holding ltd. said\nusers of its search-engine service may now access technology\nsupplied by sohu.com inc. , as the two chinese companies \nstrengthen collaboration to challenge industry leader baidu inc. users of alibaba's etao.com search service may now o...
75532360	billionaire carlos slim , the world's\nrichest man, said china must buy more and the u.s. needs to step\nup private investment to reduce the trade imbalance and boost\ntheir economies. global currency devaluation efforts will fail in the\nabsence of economic policies that f...
10176588	energy and industrial companies will\nrise next year, propelling a 17 percent gain in the standard &\npoor's 500 index from its current level, according to cambiar\ninvestors llc's brian barish . next year will be marked by a "multi-speed recovery" as\nindustries weakened b...

	path \
id	
95357231	./2008_2012_msft_bloomberg_news/netflix-quarterly-profit-increases-44-as-movie-rental-service-adds-users.txt
75227517	./2008_2012_msft_bloomberg_news/republican-sweep-may-mean-tax-boon-for-u-s-multinationals-high-incomes.txt
57850804	./2008_2012_msft_bloomberg_news/alibaba-says-it-now-offers-sohu-s-search-engine.txt
75532360	./2008_2012_msft_bloomberg_news/slim-solution-for-trade-imbalances-is-more-buying-by-china-u-s-investing.txt
10176588	./2008_2012_msft_bloomberg_news/s-p-500-to-defy-pimco-s-new-normal-rise-17-by-end-of-2011-barish-says.txt

	Month	Day	Year	\
id				
95357231	April	21	2010	

75227517	November	3	2010
57850804	October	29	2010
75532360	October	31	2010
10176588	December	1	2010

	first_sentence
id	
95357231	\n "if we had offered a pay-per-view service for new\nreleases, we would be in conflict with a broad range of\ncompanies, including wal-mart, microsoft, sony and apple," \nhastings said.
75227517	
	, microsoft corp.
57850804	
	com search service may now opt for\nthe sogou software, in addition to microsoft corp.
75532360	
	" gates, the co-founder of microsoft corp.
10176588	
	intel, microsoft corp.

Sentiment Analysis Model Application determine the sentiment of the first sentence that mentions “microsoft” or “msft” for each article.

```
[29]: def sentiment_analysis_lite(lst):
        from ds100_utils import fetch_and_cache
        import pickle
        url = "https://github.com/DS-100/sp24/raw/main/resources/assets/datasets/
        ↪hw3_sentiments.pkl"
        path = fetch_and_cache(url, "tmp-sentiment.pkl")
        with open(path, "rb") as f:
            out = pickle.load(f)
        return out
```

```
[30]: msft_news_2010['sentence_sentiment'] =_
        ↪sentiment_analysis_lite(msft_news_2010['content'].tolist())
msft_news_2010['sentence_sentiment'] = msft_news_2010['sentence_sentiment'].
        ↪map(lambda x: x['score'])
msft_news_2010.head(1)
```

Using version already downloaded: Tue May 21 20:27:18 2024

MD5 hash of file: 87a3af05678bd284866a4623c22e7d0b

```
[30]: title \
id
95357231 Netflix Profit Jumps 44% on New Users

released_at \
```

```

id
95357231  <date>April 21 2010</date> <time>23:52:36</time>

                                content  \

id
95357231  netflix inc. said first-quarter\nprofit rose 44 percent as the movie
subscription service signed\nup new customers and increased online offerings.
\n net income advanced to $32.3 million, or 59 cents a share,\nfrom $22.4
million, or 37 cents, a year earlier, the los gatos,\n...

                                path  \

id
95357231  ./2008_2012_msft_bloomberg_news/netflix-quarterly-profit-
increases-44-as-movie-rental-service-adds-users.txt

                                Month Day Year  \

id
95357231  April    21    2010

                                first_sentence  \

id
95357231  \n "if we had offered a pay-per-view service for new\nreleases, we
would be in conflict with a broad range of\ncompanies, including wal-mart,
microsoft, sony and apple,"\nhastings said.

                                sentence_sentiment

id
95357231                                0.998621

```

1.5.2 Additional Dataset

turn to an alternative, more accurate way of determining the sentiment score of articles — getting the sentiment based on the entire text, rather than getting sentiment based on the first sentence including “microsoft” or “msft” in the text. load new which contains sentiment scores of the full articles as a `DataFrame` `full_sentiments`. In this file, provided with logs which include the `id`, `score`, and `label` (“N” for “NEGATIVE” and “P” for “POSITIVE”) in the following format:

```

<device:1> <id:77243971> <result: [0.9963290095329285 (N)]>
<device:0> <id:14799046> <result: [0.9980687499046326 (N)]>
<device:1> <id:43064156> <result: [0.997868537902832 (N)]>
<device:1> <id:29402508> <result: [0.9924335479736328 (N)]>
...

```

```
[31]: full_sentiments = pd.read_csv('data/article_sentiment_logs.csv')
      full_sentiments.head()
```

```
[31]: RunNum                                log
      0      0 <device:0> <id:77243971> <result: [0.9963290095329285 (N)]>
```

```

1      1 <device:0> <id:14799046> <result: [0.9980687499046326 (N)]>
2      2 <device:0> <id:43064156> <result: [0.997868537902832 (N)]>
3      3 <device:0> <id:29402508> <result: [0.9924335479736328 (N)]>
4      4 <device:0> <id:71427879> <result: [0.9897157549858093 (N)]>

```

STEPS:

Using the logs, modify `full_sentiments` so it ultimately just contains the `id` and `content_score` (a number ranging from -1 to 1).

Then, merge this with `msft_news_2010` so we can see the results of our two methods of calculating sentiment side by side.

Assign this merged `DataFrame` to `msft_scores_2010`.

After the merge, make sure that only articles from 2010 appear and that the index of the `DataFrame` is the article `id`.

need to negate the score of negatively classified articles (indicated by “N”).

The articles have a primary key `id`.

reference how you calculated sentiment scores above.

```

[32]: regpattern = '<id:(\d+)'
      regpattern_score = '\[(.*\)]'
      updated_scores_df = pd.DataFrame(columns = ['id', 'raw_score', 'letter', 'content_score'])
      updated_scores_df['id'] = full_sentiments['log'].str.extract(regpattern).astype('int')
      updated_scores_df['raw_score'] = full_sentiments['log'].str.extract(regpattern_score)
      updated_scores_df['letter'] = updated_scores_df['raw_score'].str.split().str[1]
      updated_scores_df['raw_score'] = updated_scores_df['raw_score'].str.split().str[0].astype('float')
      updated_scores_df['content_score'] = updated_scores_df.loc[updated_scores_df['letter'] == '(N)']['raw_score'] * -1
      updated_scores_df['content_score'] = updated_scores_df['content_score'].fillna(value=updated_scores_df['raw_score'])
      updated_scores_df = updated_scores_df.loc[:, ['id', 'content_score']].set_index('id')
      msft_scores_2010 = updated_scores_df.merge(msft_news_2010, left_index = True, right_index = True)
      msft_scores_2010.head(1)

```

```

[32]: content_score      title \
id
78274823      -0.992596  Stocks, Euro Gain After Greece Asks for Aid

                                released_at \

```

```

id
78274823  <date>April 23 2010</date> <time>10:24:06</time>

                                content  \

id
78274823  stocks in europe rose and the euro\nstrengthened from a one-year low
against the dollar after greece\nsaid it will ask the european union for a
rescue package and\ngerman business confidence and earnings at companies
including\nvolvo ab beat forecasts. greek stocks and bonds ...

    path  \
id
78274823  ./2008_2012_msft_bloomberg_news/stocks-euro-gain-after-greece-asks-
for-aid.txt

    Month Day Year                                first_sentence  \
id
78274823  April   23   2010  gains were\nlimited as microsoft corp.

    sentence_sentiment
id
78274823              0.996355

```

1.5.3 Deeper Dive

```

[35]: msft_scores_2010['sentiment_difference'] =_
      ↪msft_scores_2010['sentence_sentiment'] - msft_scores_2010['content_score']
msft_scores_2010.head(1)

```

```

[35]:          content_score                                title  \
id
78274823      -0.992596  Stocks, Euro Gain After Greece Asks for Aid

                                released_at  \
id
78274823  <date>April 23 2010</date> <time>10:24:06</time>

                                content  \

id
78274823  stocks in europe rose and the euro\nstrengthened from a one-year low
against the dollar after greece\nsaid it will ask the european union for a
rescue package and\ngerman business confidence and earnings at companies
including\nvolvo ab beat forecasts. greek stocks and bonds ...

    path  \
id
78274823  ./2008_2012_msft_bloomberg_news/stocks-euro-gain-after-greece-asks-

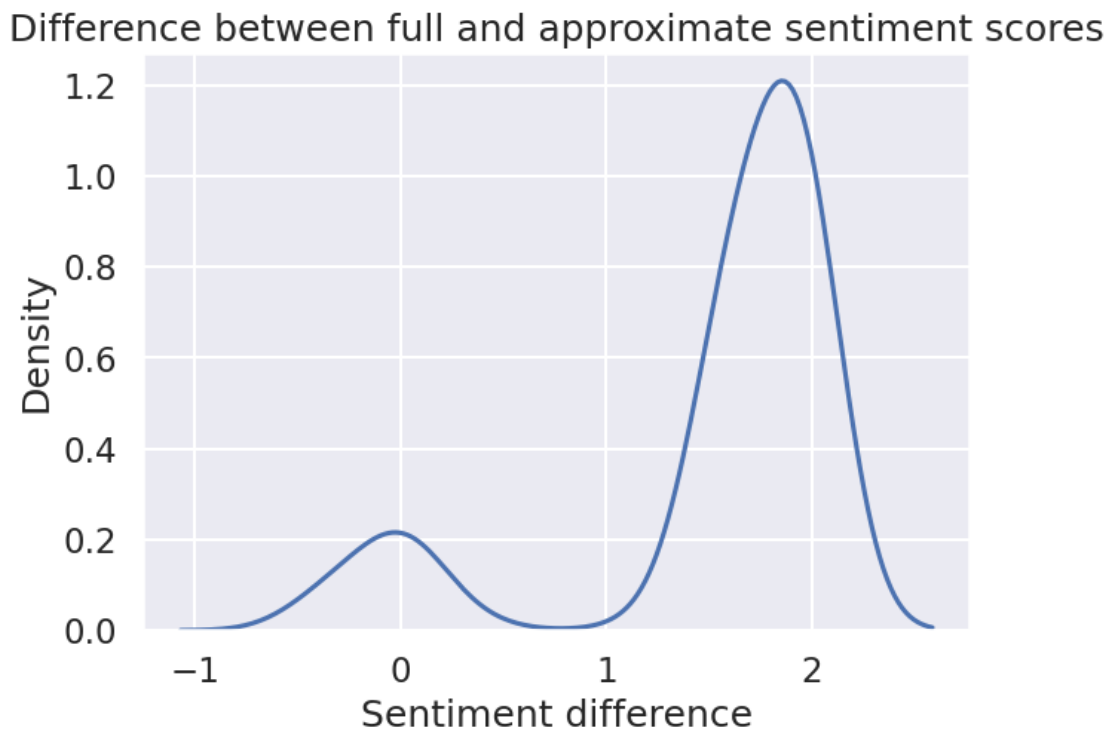
```

```
for-aid.txt
```

	Month	Day	Year	first_sentence	\
id					
78274823	April	23	2010	gains were\nlimited as microsoft corp.	

	sentence_sentiment	sentiment_difference
id		
78274823	0.996355	1.988951

```
[36]: sns.kdeplot(msft_scores_2010['sentiment_difference'])  
plt.xlabel('Sentiment difference')  
plt.title('Difference between full and approximate sentiment scores');
```



Since we were strictly observing the first sentence of the article's content, the sentiment of the entire article and its stance was unclear. Oftentimes, articles present counterarguments that can steer the sentiment detection model towards the polar opposite or an author can simply be warming up into his main argument in revealing a negative sentiment despite the first sentence of the respective article expressing neutral sentiment. Therefore, it is extremely important to gain context of the article and assess its sentiment upon full review to prevent misleading computational results. In the plot, we observe the density plot peaking at the difference of -2 with a density of >1 , highlighting the differences amongst the full article's sentiment against the first sentence's sentiment.